

Recon

Enter any URL. Get intelligence.

*Not a chatbot with a scraper.
A multi-agent recon mesh that gets smarter
with every query.*

Bright Data AI Agents Web Data Hackathon

lablab.ai | May 25-31, 2026

Team: Octo (AI Agent) + Kobus (Human)

jiwentzel@icloud.com

Competitive intelligence today is:

- **Slow:**
Manual research takes 2+ hours per company
- **Expensive:**
Semrush (\$119/mo) + ZoomInfo (\$199/mo) = \$318/mo for fragmented data
- **Fragmented:**
Analysts toggle between 5+ tools to assemble one report
- **Not scalable:**
Growth teams need intel on 50+ competitors. At 2 hours each, that's 100 hours of work.

The real cost:

While your analyst researches, your competitor ships.

Recon: Enter any company URL.

Get an analyst-grade competitive intelligence report in 10 seconds.

8-10 second

report generation

\$2.00

per report

Zero manual work

full automation

Gets smarter

persistent memory

What You Get:

- Company overview (tech stack, traffic, social presence)
- Recent news & press mentions
- Competitive positioning
- Key executives & team size
- Funding & growth indicators
- Traffic sources & SEO insights

Replaces:

Semrush + ZoomInfo + SimilarWeb + manual analyst research

→ **One \$2 report**

1. Sign Up

User Flow (3 Steps)

- Get \$10 credits
- No credit card needed

2. Enter Company URL

- Paste target URL
- Click 'Generate Report'

3. Get Report

- Live pipeline fires
- Report ready in ~8-10s
- \$2.00 debited

What Makes It Different

- **Real-time visibility:**
Watch each agent work in live waterfall view
- **Cost transparency:**
See exactly what each step costs
- **Instant caching:**
Query same domain again? Instant load from memory
- **No subscriptions:**
Pay per report, not per month

While other tools show 'Loading...', Recon shows the intelligence operation:

[007-bot]	Task received	0.0s	●	
[Circus]	Routing to BD worker	0.1s	●	
[Web Unlocker]	Scraping target URL	0.3s	■■■■■■■■■	1.4s
[SERP API]	Searching company + news	1.8s	■■■■■	0.6s
[Scrp Browser]	JS sub-pages (LinkedIn/CB)	2.5s	■■■■■■■■■	2.1s
[Web Scraper]	Extracting structured data	4.6s	■■■	0.6s
[AI-IQ]	Storing 12 facts to memory	5.2s	■	0.2s
[Claude]	Synthesizing report	5.4s	■■■■■■■■■	3.0s
[Circus]	Publishing to shared pool	8.4s	■	0.1s
[007-bot]	Report ready	8.5s	✓	

Post-delivery message:

"This intelligence is now available to all 6 mesh agents.
Next query on this domain loads from memory instantly."

06 Bright Data Integration — All 4 Products

Product	Purpose	Cost	Why Critical
Web Unlocker	Bypass bot detection on target company website	\$0.30	Most company sites block scrapers
SERP API	News & competitors search results	\$0.50	Real-time news data from Google/Bing
Scraping Browser	JS-heavy pages (LinkedIn, Crunchbase)	\$0.80	80% of B2B data behind JavaScript
Web Scraper API	Structured data extraction	\$0.40	Turns HTML into clean JSON

Report Generated: \$2.00

Breakdown:


- Web Unlocker \$0.30
 - SERP API \$0.50
 - Scraping Browser \$0.80
 - Web Scraper API \$0.40
-

Total \$2.00

Bypassed bot detection on 3 pages
Saved ~15,000 tokens vs manual browsing
Report generated in 8.4s

07 Multi-Agent Architecture

Not one chatbot with a scraper — a 6-agent intelligence mesh coordinated by Circus



Why 6 Agents > 1 Bot

- Specialized agents per task vs single bot doing everything
- Credentials isolated to BD Worker vs API keys hardcoded
- AI-IQ persists all facts vs no memory between queries
- Parallel agent execution vs serial processing
- Agent failures isolated vs one failure = total failure
- Live waterfall shows all agents vs no visibility

Problem:

Traditional scrapers refetch data every time. Wasteful and slow.

Solution:

AI-IQ persistent memory system

1. First query on stripe.com

- BD Worker scrapes all sources
- AI-IQ extracts facts: ["Stripe processes \$640B annually", ...]
- Facts stored with tags: company:stripe, category:financials
- Report generated in 8.4s

2. Second query on stripe.com

- AI-IQ checks: "Do I have recent facts on stripe.com?"
- Finds 12 cached facts from 2 minutes ago
- BD Worker skips (unless facts are >7 days old)
- Report generated in 0.3s (28× faster)

3. Third query on stripe.com/pricing

- AI-IQ returns cached company facts
- BD Worker only scrapes /pricing (incremental update)
- Report includes both cached + new data
- Report generated in 2.1s

Why This Matters for Hackathon

- **Originality:** No other entry will have persistent cross-query memory
- **Business Value:** 10 users query same company, Recon pays Bright Data once, not 10×
- **Tech Depth:** Shows production concerns (caching, cost optimization)

Pricing: Pay-Per-Report

Mode	Scouts	Time	Cost	Use Case
Standard	4 BD products	~9s	\$2.00	Quick company overview
SEO	6 scouts	~18s	\$5.00	Search presence audit
Red Team	8 scouts (OSINT)	~30s	\$12.00	Attack surface mapping
Deep Search	10 scouts	~24s	\$15.00	Full competitive intel
Bundle	All of above	~45s	\$25.00	Complete intelligence

Comparison: Recon vs Incumbents

Solution	Cost	Time	Depth
Manual Analyst	\$75/hr × 2hr = \$150	2 hours	High
Semrush + ZoomInfo	\$318/mo (÷30 = \$10.60/day)	30 min + manual	Medium (fragmented)
Recon Standard	\$2.00	10 seconds	High (unified)
Recon Bundle	\$25.00	45 seconds	Analyst + SEO + Security

Target Markets

• Growth Teams (primary)

- Need intel on 50+ competitors quarterly
- 50 reports $\times$ \$2 = \$100 vs Semrush (\$119/mo $\times$ 12 = \$1,428/yr)

• Investors / VCs

- Due diligence on 200+ companies per year
- 200 reports $\times$ \$2 = \$400 vs ZoomInfo (\$199/mo $\times$ 12 = \$2,388/yr)

• Recruiters

- Research target companies before outreach
- 100 reports $\times$ \$2 = \$200 vs LinkedIn Sales Nav (\$79/mo $\times$ 12 = \$948/yr)

• Journalists

- Background research for articles
- 30 reports $\times$ \$2 = \$60 vs Semrush (\$119/mo = \$119)

TAM (Total Addressable Market)

- Competitive intelligence market: \$8.2B by 2025 (Verified Market Research)
- Web scraping services market: \$1.8B by 2027 (Market Research Future)
- Overlap (our niche): ~\$500M = companies that need intelligence but can't afford analysts

\$500M Addressable Market

11 Originality: What No Other Entry Has

1. Live Multi-Agent Waterfall Visualization

Real-time progress bars showing each agent's work. Every other tool shows 'Loading...' — we show the intelligence operation.

2. Persistent Cross-Query Memory

AI-IQ stores facts across sessions. Second query on same domain = instant (cached).

3. Cost Transparency

User sees exactly what each Bright Data product cost per report. No 'contact sales' hidden pricing.

4. Agent Delegation Chain

007-bot → Circus → BD Worker → AI-IQ → Reporter. Not just 'call API in sequence' — proper mesh routing.

5. SaaS-Quality UI (Not a Demo)

Credits system, live waterfall panel, downloadable reports, cost breakdown. Feels like a product you'd pay for today.

Expected Competition


Most hackathon entries will be:

- Single-agent chatbots that call Bright Data API
- Show scraped data in JSON or table
- No memory, no caching, no cost breakdown
- No waterfall visualization

Recon is the only entry that:

- ✓ Uses all 4 Bright Data products in one pipeline
- ✓ Shows live agent topology
- ✓ Has persistent memory across queries
- ✓ Displays cost transparency
- ✓ Feels like a real SaaS product

Side-by-Side Panels: Intelligence Pipeline + Report



Key UI Elements:

- Live Waterfall (left) — Shows each agent, progress bars animate, checkmarks when complete
- Report Panel (right) — Markdown-formatted, download/share buttons, cost breakdown
- Credits Display (top) — Shows remaining balance, updates in real time
- Post-Report Message — 'This intelligence is now in AI-IQ memory. Next query instant.'

Frontend

- Next.js / React (UI)
- Server-Sent Events (SSE) for live waterfall updates
- TailwindCSS (styling)

Backend

- Circus mesh (agent coordination)
- 007-bot (entry point, Node.js)
- BD Worker (Bright Data orchestrator, Node.js)
- AI-IQ (memory system, SQLite + vector embeddings)
- Reporter (Claude API for synthesis)

Bright Data Products

- Web Unlocker (bypass bot detection)
- SERP API (news & competitors)
- Scraping Browser (JS-heavy pages)
- Web Scraper API (structured extraction)

Infrastructure

- VPS (WhatsApp server)
- Circus running on localhost:6200
- Agent workspaces: /root/octo-workspace
- Memory DB: /root/.ai-iq/meshint.db

Key Libraries

- playwright (fallback scraping)
- cheerio (HTML parsing)
- marked (markdown rendering)
- memory-tool (AI-IQ CLI)

Week 1-2: Polish & Launch

- Final branding
- Payment integration (Stripe for credits)
- User authentication (OAuth + magic links)
- Public launch on Product Hunt

Month 1: Feature Expansion

- Bulk reports: Upload CSV of 50 URLs
- Scheduled monitoring: Alert when competitor gets news
- Comparison mode: Side-by-side reports for 2-5 companies
- Export formats: PDF, Google Docs, Notion integration

Month 2-3: Team Plans

- Shared workspaces: Growth team shares report history
- API access: Integrate Recon into internal tools
- Slack bot: /meshint stripe.com = report in channel
- Custom agents: Upload your own scrapers to mesh

Month 4-6: Intelligence Layer

- Trend detection: '5 SaaS companies raised Series A this week'
- Anomaly alerts: 'competitor.com traffic dropped 40%'
- Predictive insights based on hiring patterns
- Knowledge graph: Visualize company connections

Long-Term Vision

Become the operating system for competitive intelligence
Not just reports — real-time alerts, trend analysis, predictive insights

Target: 10,000 reports/day within 6 months

Octo (AI Agent)

Role: Lead developer, architecture, implementation

Agent ID: octo-7aea1b

Workspace: /root/octo-workspace

Capabilities: Full-stack development, agent coordination, memory systems

Built: Circus mesh integration, AI-IQ memory layer, waterfall visualization

Kobus (Human)

Role: Product strategy, Bright Data integration, pitch

Email: jiwentzel@icloud.com

Background: Builder of WhatsHub (agent mesh infrastructure)

Responsibilities: Bright Data API setup, hackathon submission, demo recording

Why This Team Works

- Octo handles implementation complexity (6 agents, live updates, memory)
- Kobus handles product clarity (what users see, why they pay)
- Result: Technical depth + business polish

Built at the intersection of AI agents and competitive intelligence

Enter any URL.
Get intelligence.
In 10 seconds.

*"Every AI agent eventually hits a wall.
We built the infrastructure to go through it."*

Try with any company URL:

stripe.com
openai.com
anthropic.com
your-competitor.com

Contact: jiwentzel@icloud.com
Built for Bright Data AI Agents Web Data Hackathon
May 25-31, 2026 | lablab.ai

Where to Find Each Judging Criterion

• Application of Technology

- Cost breakdown shows all 4 BD products used
- Waterfall visualization shows them firing in parallel
- Every report uses Web Unlocker, SERP API, Scraping Browser, Web Scraper API

• Presentation

- UI feels like a \$20/mo SaaS product
- Credits system, live updates, downloadable reports
- Professional polish, not hackathon demo quality

• Business Value

- \$2/report vs \$318/mo subscriptions
- Replaces 2 hours of analyst time with 10-second automation
- Credits system proves payment model works

• Originality

- Live agent waterfall — no other entry has this
- Persistent memory across queries — cost optimization
- Cost transparency per Bright Data product
- 6-agent mesh architecture vs single-bot approach

Why Recon Wins

Uses all 4 Bright Data products • Shows real business value • Feels like a real product
Technical originality with multi-agent mesh • Clear post-hackathon path

Not a chatbot with a scraper. A multi-agent recon mesh.